

---

--- title: "Tools & Concepts for Cloud Deployments" author: ["Christopher B. Hauser"] institute: ["Institute of Information Resource Management, Ulm University"] subject: "Markdown" tags: [Markdown, Example] titlepage: true graphics: true mainfont: Open Sans mainfontoptions: BoldFont=Open Sans Bold mainfontoptions: ItalicFont=Open Sans Italic mainfontoptions: BoldItalicFont=Open Sans Bold Italic

date: SummerSchool 2019, Curitiba subtitle: "Solution for Exercise 2" --- # Lesson 1: Monitoring with InfluxData

## How your Setup should look like

You should have two instances in Openstack:

![Instances in OpenStack](imgs/openstack-instances.png)

The Chronograf dashboard for your main-server instance should look like follows:

![Chronograf Dashboard](imgs/chronograf.png)

## Question: Parts and responsibilities of the TICK stack

![Components of InfluxData TICK stack](imgs/tickstack.png)

## Question: What ports are used by influxdb? By default, InfluxDB uses the following network ports: - TCP port 8086 is used for client-server communication over InfluxDB's HTTP API - TCP port 8088 is used for the RPC service for backup and restore

“ ubuntu@monitoring:~\$ sudo netstat -tulpen sudo: unable to resolve host monitoring Active Internet connections (only servers) Proto Recv-Q Send-Q Local Address Foreign Address State User Inode PID/Program name tcp 0 0 0.0.0.0:22 0.0.0.0:\* LISTEN 0 13017 1043/sshd tcp6 0 0 :::8086 :::\* LISTEN 999 18068 2411/influxd tcp6 0 0 :::22 :::\* LISTEN 0 13025 1043/sshd tcp6 0 0 :::8088 :::\* LISTEN 998 18308 2505/chronograf tcp6 0 0 :::8088 :::\* LISTEN 999 18066 2411/influxd udp 0 0 0.0.0.0:68 0.0.0.0:\* 0 13398 880/dhclient “

# Lesson 2: Vertical Scaling

## Stressing Mediawiki with small flavor

With a single thread on client side:

“ Concurrency Level: 1 <<<< Time taken for tests: 618.219 seconds Complete requests: 5000 Failed requests: 4958 (Connect: 0, Receive: 0, Length: 4958, Exceptions: 0) Non-2xx responses: 5000 Total transferred: 58650042 bytes HTML transferred: 56685042 bytes Requests per second: 8.09 [# /sec] (mean) Time per request: 123.644 [ms] (mean) Time per request: 123.644 [ms] (mean, across all concurrent requests) Transfer rate: 92.65 [Kbytes/sec] received

Connection Times (ms) min mean[+/-sd] median max Connect: 0 0 0.8 0 12 Processing: 101 123 21.5 121 1124 Waiting: 55 71 14.2 68 813 Total: 101 124 21.5 121 1124

---

Percentage of the requests served within a certain time (ms) 50% 121 66% 125 75% 127 80% 129 90% 135 95% 142 98% 152 99% 162 100% 1124 (longest request) ““(imgs/bench-single.png)

With ten concurrent client side threads:

““ Concurrency Level: 10 <<<< Time taken for tests: 165.952 seconds Complete requests: 5000 Failed requests: 4 (Connect: 0, Receive: 0, Length: 4, Exceptions: 0) Non-2xx responses: 5000 Total transferred: 58654996 bytes HTML transferred: 56689996 bytes Requests per second: 30.13 [# /sec] (mean) Time per request: 331.903 [ms] (mean) Time per request: 33.190 [ms] (mean, across all concurrent requests) Transfer rate: 345.16 [Kbytes/sec] received

Connection Times (ms) min mean[+/-sd] median max Connect: 0 0 0.5 0 11 Processing: 89 332 57.8 333 501 Waiting: 76 284 50.9 285 442 Total: 90 332 57.8 333 501

Percentage of the requests served within a certain time (ms) 50% 333 66% 357 75% 372 80% 380 90% 406 95% 425 98% 447 99% 459 100% 501 (longest request) ““(imgs/bench-10concurrent.png)

## Stressing Mediawiki after vertical scaling

With a single thread on client side:

““ Concurrency Level: 1 Time taken for tests: 604.299 seconds Complete requests: 5000 Failed requests: 4904 (Connect: 0, Receive: 0, Length: 4904, Exceptions: 0) Non-2xx responses: 5000 Total transferred: 58650096 bytes HTML transferred: 56685096 bytes Requests per second: 8.27 [# /sec] (mean) Time per request: 120.860 [ms] (mean) Time per request: 120.860 [ms] (mean, across all concurrent requests) Transfer rate: 94.78 [Kbytes/sec] received

Connection Times (ms) min mean[+/-sd] median max Connect: 0 0 0.1 0 2 Processing: 101 121 20.6 117 915 Waiting: 53 69 19.6 65 801 Total: 101 121 20.6 117 916

Percentage of the requests served within a certain time (ms) 50% 117 66% 121 75% 124 80% 126 90% 132 95% 140 98% 159 99% 177 100% 916 (longest request) ““

!(imgs/bench-vscaled-single.png)

With ten concurrent client side threads:

““ Concurrency Level: 10 Time taken for tests: 79.773 seconds Complete requests: 5000 Failed requests: 1324 (Connect: 0, Receive: 0, Length: 1324, Exceptions: 0) Non-2xx responses: 5000 Total transferred: 58653676 bytes HTML transferred: 56688676 bytes Requests per second: 62.68 [# /sec] (mean) Time per request: 159.547 [ms] (mean) Time per request: 15.955 [ms] (mean, across all concurrent requests) Transfer rate: 718.02 [Kbytes/sec] received

Connection Times (ms) min mean[+/-sd] median max Connect: 0 0 0.1 0 2 Processing: 61 159 36.8 156 313 Waiting: 53 134 28.8 132 272 Total: 61 159 36.8 156 314

---

Percentage of the requests served within a certain time (ms) 50% 156 66% 171 75% 182 80%  
190 90% 209 95% 225 98% 243 99% 258 100% 314 (longest request) ““



## Conclusion

CPU cores	Threads	Requests / s	---	---	---	2	1	8 Req/s	2	10	30 Req/s	4
1	8 Req/s	4	10	60 Req/s								

- Application is CPU bound - For concurrent requests more vCPUs are improving a lot! - Virtual scaling is limited to the largest available flavor. - => unlimited scalability is not vertical.