



Tools & Concepts for Cloud Deployments

Exercise 6: Container Orchestration 1

Christopher B. Hauser

Institute of Information Resource Management, Ulm
University

Exercise 6: Container Orchestration 1

Overview

Now that we have learned how to work with containers on a single node, it is time to orchestrate Docker on multiple hosts. This exercise will start with:

- How to install and use Docker Swarm
- How to install and use Rancher
- What are core features of Container Orchestrators
- How to deploy and use a private Docker Registry

Lessons

1. Container Orchestration with Docker Swarm
2. Container Orchestration with Rancher

[Exercise as PDF](#)

Exercise Solution

- [Solution as Markdown](#)
- [Solution as PDF](#)

Lesson 1: Container Orchestration with Docker Swarm

In this lesson, we will install and use Docker Swarm, as the built in Docker orchestrator across multiple hosts.

Research: How Docker Swarm works

Make yourself familiar with Docker Swarm. The two links should give you enough information to work with Swarm in the following:

<https://docs.docker.com/engine/swarm/how-swarm-mode-works/nodes/>

<https://docs.docker.com/engine/swarm/how-swarm-mode-works/services/>

Questions: Docker Swarm

- In the Swarm terminology, what are services, tasks, and containers?
- Where in our Cloud Stack do you place Docker Swarm?

Cloud Stack	Example	Deployment Tool
Application Component	Mediawiki	?
Containers	Docker	?
Virtual Resource	Instance m1.small	?
Cloud Platform	OpenStack	-

Task: Setup Docker Swarm

Deploy three virtual machines with docker installed on OpenStack. We can use the terraform bundle "terraform-dockernodes.zip" to automatise this process. Update the `provider.tf` as we did in exercise 4.

Log in to the three VMs via SSH. Validate that Docker is running (e.g. via `docker ps -a`). Docker Swarm is inactive, validate output of `docker info` and search for line `Swarm: inactive`.

Note: if you get `docker: Got permission denied while trying to connect to the Docker daemon socket` the cloud-init script has not yet finished. Logout and login a few seconds later again.

Lets take `dockernode1` as swarm manager. Log in to `dockernode1`, look up the private IP address of this VM (`$PRIV_IP_DOCKERNODE1`), then run `docker swarm init --advertise-addr`

`$PRIV_IP_DOCKERNODE1`. The output of this command returns:

Swarm initialized: current node (....) is now a manager.

To add a worker to this swarm, run the following command:

```
docker swarm join \
  --token SWMTKN-1..... \
  192.168.5.4:2377
```

Copy/Paste the `docker swarm join` command to `dockernode2` and `dockernode3`. Validate your Docker Swarm cluster on `dockernode1` with `docker node ls` - you should see three hosts.

Task: Start a test service

For testing purposes we will now start three instances of the Ghost blog software on our Swarm cluster. Note: `docker service ...` commands can only be executed on the swarm manager, in our case `dockernode1`. Create the service with 3 replicas:

```
docker service create --replicas 3 --name ghost-blog \
  --publish 80:2368 ghost
```

Validate that the service is booting with `docker service ls`, where the counter in column "REPLICAS" will go from 0/3 to 3/3 eventually. You can see that the three replicas are distributed among the three nodes: `docker service ps ghost-blog`. On each virtual machine run `docker ps`, you will have one container of the ghost service on each node. Take a browser, and go to any of your three floating IPs on port 80 to access the ghost instances.

Swarm allows you to scale horizontally up and down (`docker service scale $SERVICE_NAME=$NUMBER_OF_REPLICAS`). You can easily scale down the ghost service to only one replica with `docker service scale ghost-blog=1`. Validate the number of replicas with `docker service ps ghost-blog`. Note: if your container runs on e.g. `dockernode1` you can access the ghost blog via `http://$FLOATING_IP_DOCKERNODE1/`. But Swarm handles published ports across all swarm nodes. You can also access the ghost blog via any other floating ip.

Updating the software is also handled by Swarm, where sequentially containers are removed and recreated to achieve zero downtime of the service. In case of errors, Swarm can rollback to the previous version.

Remove the test service with `docker service rm ghost-blog` before you continue.

Task: Mediawiki with Docker Swarm

Docker Swarm and Docker Compose work together (software is experimental at present).

So far we built the images from Dockerfiles on the node where we started containers from those images. Since we're having a distributed Swarm cluster, we need a central store for our built images: a (private) docker registry. We could use Docker Hub, or deploy our own registry. On dockernode1, run the following commands to run your own docker registry:

```
# install certbot to get certificate
sudo apt-get install -y software-properties-common
sudo add-apt-repository -y ppa:certbot/certbot
sudo apt-get update
sudo apt-get install -y --allow-unauthenticated certbot

# lookup your public hostname
floatingIP=$(curl -s http://169.254.169.254/latest/meta-data/public-
ipv4)
publicHostname=$(nslookup $floatingIP | grep "in-addr.arpa" |\
    cut -d'=' -f2 | tr -d ' ')

# get certificate
sudo certbot certonly --standalone -d $publicHostname --agree-tos \
    --register-unsafely-without-email --preferred-challenges http
sudo mkdir /opt/certs
folder=$(sudo ls /etc/letsencrypt/live/)
sudo cp /etc/letsencrypt/live/$folder/fullchain.pem /opt/certs/
sudo cp /etc/letsencrypt/live/$folder/privkey.pem /opt/certs/

# finally: start your docker registry as container
docker run -d -p 5000:5000 --restart=always --name registry \
    -v /opt/certs:/certs \
    -e REGISTRY_HTTP_TLS_CERTIFICATE=/certs/fullchain.pem \
    -e REGISTRY_HTTP_TLS_KEY=/certs/privkey.pem \
    registry:2
```

Next, copy the Dockerfiles from exercise 5 to dockernode1. Change the `docker-compose.yml` file to the following content (replace `omistack-vmXYZ.e-technik.uni-ulm.de` to fit your `$publicHostname`):

```
version: '3'
services:
  web:
    build: Mediawiki
```

```
    image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/mediawiki
  ports:
    - 80:80
  deploy:
    replicas: 2
    mode: replicated
database:
  build: Database
  image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/database
  deploy:
    replicas: 1
    mode: replicated
```

Then build and push the images to your registry (still on dockernode1):

```
docker-compose build
docker push omistack-vmXYZ.e-technik.uni-ulm.de:5000/database
docker push omistack-vmXYZ.e-technik.uni-ulm.de:5000/mediawiki
```

Validate, that the registry works. Go to dockernode2 and download the image via `docker pull omistack-vmXYZ.e-technik.uni-ulm.de:5000/database`

Now we can use Docker Swarm to deploy our Docker Compose file. On dockernode1, run:

```
docker stack deploy --compose-file docker-compose.yml mediawiki
```

Validate with `docker stack ps mediawiki` that you have three containers running, and check if you can access the mediawiki application via web browser.

We have now deployed the mediawiki in Docker containers across several (manually created) Docker hosts.

Task: Clean up

Depending on your interests, run experiments with Docker Swarm. When you're done, make sure to clean up the resources in OpenStack, especially before moving on to the next lesson.

Clean up with terraform is as easy as typing `terraform destroy :-)`

Lesson 2: Container Orchestration with Rancher

Docker Swarm offers clustering of hosts to one large container cluster, and orchestrates services by placing containers on available hosts. Rancher is very equal, yet it even offers a web ui to manage a Rancher container cluster.

Research: Rancher

What is Rancher? Become familiar with the tool before we start.

<http://rancher.com/>

<http://docs.rancher.com/rancher/v1.6/en/>

Question: Rancher

Where in our Cloud Stack do you place Rancher?

Cloud Stack	Example	Deployment Tool
Application Component	Mediawiki	?
Containers	Docker	?
Virtual Resource	Instance m1.small	?
Cloud Platform	OpenStack	-

Task: Install and start Rancher

Connect via ssh to your rancher-master VM. Now start a Rancher service (it takes several seconds until rancher is accessible):

```
docker run -d --restart=unless-stopped -p 8080:8080 --name rancher \
  rancher/server:stable
```

Note: if you get `docker: Got permission denied while trying to connect to the Docker daemon socket: the cloud-init script has not yet finished`. Logout and login a few seconds later again.

Next, we will add hosts (virtual machines with a Rancher agent) to Rancher. Open the web dashboard of your Rancher Master http://PUBLIC_IP:8080/.

- Click on "Infrastructure > Hosts" and select "Add Host".
- In "Host Registration URL" select "something else" and type in `http://PRIVATE_IP:8080` with the private ip of your rancher-master. Press Save.
- On the next site, copy the `sudo docker ...` command from step 5 and execute it on all your rancher-host VMs and also to your rancher-master.
- Click "Close" and wait for the hosts to show up in the rancher dashboard.

Starting the rancher hosts will take some time. You can follow the process: Rancher agents will be started, and several predefined containers like healthchecker, or scheduler will be installed. When the hosts are up, you can browse through the utilisation metrics and the deployed containers in the web dashboard. When selecting single containers, you get statistics about this container only, its even possible to get execution shells into containers from the dashboard.

Finally install the Rancher CLI tool inside your initial rancher VM:

```
wget -O rancher.tar.gz \
    https://github.com/rancher/cli/releases/download/v0.6.1/rancher-
linux-amd64-v0.6.1.tar.gz
tar xfv rancher.tar.gz
sudo mv rancher-v0.6.1/rancher /usr/bin/
sudo chmod +x /usr/bin/rancher
```

Besides the web dashboard you can now connect to rancher via CLI.

To authenticate in the CLI, you need an API Key. In the web dashboard, go to **API > Keys** and press the **Add Account API Key** Button. Enter any name and press create. You will get an Access Key and a Secret Key. Copy both to a place where you will find them later.

To use the CLI, run `rancher config` and provide URL `[]`: `http://134.60.64.XYZ:8080/v1`, the Access Key and Secret Key from before. Validate that the login works with `rancher ps`.

Task: Start Mediawiki via Rancher

Copy and unzip the `dockerfiles.zip` from exercise 4 to the rancher vm. Inside the extracted `dockerfiles` folder, create a `docker-compose.yml` with the following content:

```
version: '2'
services:
  web:
    build: Mediawiki
    image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/mediawiki
    ports:
      - 80:80
```



```
database:
  build: Database
  image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/database
```

Then build and push the images to your registry:

```
docker-compose build
docker push omistack-vmXYZ.e-technik.uni-ulm.de:5000/database
docker push omistack-vmXYZ.e-technik.uni-ulm.de:5000/mediawiki
```

Now we are ready to import the docker-compose file as a stack into Rancher, either by web dashboard or by cli.

In the web dashboard, go to “Stacks”, and press the “Add Stack” button. Add as name “mediawiki” and place as content of docker-compose.yml:

```
version: '2'
services:
  web:
    image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/mediawiki
    ports:
      - 80:80
  database:
    image: omistack-vmXYZ.e-technik.uni-ulm.de:5000/database
```

Mediawiki will be started. Try to access it. You will recognize, that it is only accessible via the floating ip of the host where the web container is running.

To allow any of the rancher hosts, we have to add a load balancer: In the mediawiki stack, click “add service” and select “load balancer”. Scale: Always run one instance of this container on every host Name: wiki-loadblancer Request Port 80, Target: web Port 80.

When selecting the web service, you can configure the Health Check, so rancher will restart failed containers automatically e.g. when the web server is not responding (Docker will only react on stopped containers). When you edit a service, you can scale it up and down. The loadbalancer will be configured automatically.